

# Indexers on Sui: A Modular and Generic Approach to Blockchain Data Processing

Master Project

May 20, 2025

## Contents

<b>1 Introduction</b>	<b>2</b>	3.5 Advantages over Traditional Approaches	8
1.1 The Critical Role of Indexers in Blockchain Ecosystems	3	3.6 Use Case Example: NFT Collection Analytics	8
1.2 Current Challenges in Data Indexing on Sui	3	3.7 Limitations and Future Work	8
1.3 Towards a More Accessible Indexing Solution	3	<b>4 Contribution 2: Generic Package-Based Indexer</b>	<b>9</b>
<b>2 Background and Current State of Indexing on Sui</b>	<b>4</b>	4.1 Design Philosophy	9
2.1 Sui Blockchain Architecture	4	4.2 Technical Implementation	9
2.2 Data Flow and Transaction Lifecycle in Sui	4	4.2.1 Architecture Overview	9
2.2.1 Transaction Structure	4	4.2.2 Key Implementation Features	9
2.2.2 Checkpoint Structure	5	4.3 Deployment and Usage	10
2.3 Current Indexing Methods and Schemas	5	4.3.1 Local vs. Remote Mode	10
2.3.1 Native Sui Indexing	5	4.4 Data Model and Storage	10
2.3.2 Current Third-Party Indexing Solutions	5	4.5 Use Case Example: DeFi Protocol Analytics	10
2.3.3 Typical Indexing Schema	5	4.6 Advantages for Non-Technical Teams	11
2.4 Limitations and Challenges of Current Approaches	5	4.7 Limitations and Future Work	11
2.5 The Developer Experience Gap	6	<b>5 Analysis and Comparative Evaluation</b>	<b>11</b>
<b>3 Contribution 1: Modular Indexer SDK Approach</b>	<b>6</b>	5.1 Comparative Analysis of the Two Contributions	11
3.1 Design Philosophy	6	5.2 Comparison with Existing Solutions	12
3.2 Technical Implementation	7	5.2.1 Feature Comparison	12
3.2.1 Architecture Overview	7	5.2.2 Development Time Impact	12
3.2.2 Core Components Analysis	7	5.3 Performance Analysis	12
3.3 Database Schema Design	7	5.3.1 Processing Efficiency	12
3.4 Implementation Walkthrough	7	5.3.2 Resource Utilization	12
		5.4 Developer Experience Enhancement	13
		5.5 Integration with Analytics Platforms	13
		5.6 Ecosystem Impact Assessment	13
		5.7 Future Research Directions	13

<b>6</b>	<b>Conclusion</b>	
6.1	Summary of Contributions . . . . .	14
6.2	Impact on the Developer Experience .	14
6.3	Limitations and Future Work . . . . .	15
6.4	Broader Implications for Blockchain Ecosystems . . . . .	15
6.5	Final Thoughts . . . . .	15

## 14 Abstract

This master project addresses a significant challenge in the Sui blockchain ecosystem: the time-consuming and resource-intensive process of data indexing. Our research reveals that development teams spend approximately 30% of their project time implementing indexing solutions to access and analyze on-chain data. We introduce two complementary contributions to address this bottleneck: (1) a modular SDK (corresponding to the `sui-indexer-modular`) that allows developers to specify both the database type and the precise data to extract from specific transaction fields, and (2) a so-called generic indexer (implemented as `sui_indexer_checkpointTx`) that automatically processes all transactions related to a specific package but does not allow the developer to choose the database type or customize the data extraction. Both solutions prioritize data self-custody, allowing teams to run indexers locally or in their own infrastructure. Our implementations significantly reduce the technical barrier to entry for data indexing on Sui, enabling development teams to focus on their core application logic while maintaining full control over their data. Comparative analysis with existing solutions demonstrates the efficiency and accessibility advantages of our approach, particularly for teams without specialized blockchain knowledge or dedicated indexing resources.

## 1 Introduction

In the rapidly evolving landscape of blockchain technology, the ability to efficiently access and analyze on-chain data has become a critical requirement for decentralized applications (dApps). As blockchain ecosystems mature, the volume of transaction data grows exponentially, creating significant challenges for developers who need to access this data in real-time to power their applications.

The Sui blockchain, developed by Mysten Labs, represents a new generation of high-performance blockchains designed for wide-scale adoption. With its unique object-centric data model and high throughput capabilities, Sui enables developers to

build sophisticated dApps with enhanced performance and scalability. However, this advancement comes with its own set of challenges, particularly in the realm of data indexing and retrieval.

## 1.1 The Critical Role of Indexers in Blockchain Ecosystems

Blockchain indexers play a pivotal role in the ecosystem by transforming raw on-chain data into structured, queryable formats that applications can easily consume. Without efficient indexing solutions, dApps would need to scan the entire blockchain history to retrieve relevant information, a process that is both resource-intensive and time-prohibitive.

On the Sui blockchain, indexers are especially important due to several factors:

- **Real-time Frontend Updates:** Modern dApps require near-instantaneous reflection of on-chain state changes in their user interfaces.
- **Analytics and Insights:** Teams need comprehensive data about user interactions, transaction patterns, and contract performance to optimize their applications.
- **Historical Data Access:** Applications often need to access historical transaction data that may not be readily available from full nodes.
- **Custom Filtering and Processing:** Different applications have unique requirements for filtering and processing on-chain events relevant to their specific use cases.

## 1.2 Current Challenges in Data Indexing on Sui

Despite their critical importance, implementing effective indexing solutions on Sui presents several significant challenges:

- **Technical Complexity:** Building a custom indexer requires deep knowledge of Rust programming and intimate familiarity with the Sui framework, creating a steep learning curve for many development teams.

- **Resource Requirements:** Traditional indexing approaches often require dedicated development resources, diverting attention from core application development.
- **Data Custody and Ownership:** Existing third-party solutions may not provide full data ownership or self-custody options, creating potential dependencies and restrictions.
- **Scalability Concerns:** As applications grow, their data requirements increase, requiring indexing solutions that can scale efficiently with minimal overhead.
- **Cost Efficiency:** For early-stage projects or smaller teams, the cost of developing and maintaining custom indexing infrastructure can be prohibitive.

A recent survey of Sui developers and founders, conducted in collaboration with the Sui DevRel team, revealed that implementing proper data indexing solutions consumes approximately 30% of the total development time required to bring a product to market on the Sui blockchain. This represents a significant bottleneck in the development pipeline and highlights the urgent need for more efficient and accessible indexing solutions.

## 1.3 Towards a More Accessible Indexing Solution

This master project addresses these challenges by introducing two complementary contributions to the Sui ecosystem:

1. A modular SDK approach (`sui-indexer-modular`) that allows developers to specify both the database type and precise actions on specific fields of checkpoint transactions, simplifying the framework and reducing the learning curve.
2. A so-called generic indexing solution (`sui_indexer_checkpointTx`) that automatically fetches all transaction data related to a specified package, but does not allow the developer to choose the database type or customize

the data extraction, enabling teams to run analytical operations without deep knowledge of the Sui framework.

Both solutions are designed with self-custody in mind, allowing teams to maintain full ownership of their data by running the indexer locally or in their own infrastructure. By providing these tools, we aim to significantly reduce the development overhead associated with data indexing, allowing teams to focus on their core application logic while still benefiting from comprehensive on-chain data access.

The following sections will detail the technical implementation of these solutions, explore their benefits compared to existing approaches, and demonstrate their practical application through real-world use cases.

## 2 Background and Current State of Indexing on Sui

To understand the challenges and opportunities in Sui blockchain data indexing, it is essential to first establish a clear understanding of the Sui architecture, data model, and current indexing approaches.

### 2.1 Sui Blockchain Architecture

Sui is a high-performance Layer 1 blockchain designed with an object-centric data model, which differentiates it from account-based blockchains like Ethereum. In Sui, the fundamental unit of storage is an object rather than an account, and each object has a unique identifier (ID) that persists throughout its lifetime.

Key architectural elements of Sui include:

- **Object-Centric Model:** All on-chain state is represented as objects that can be created, modified, or deleted through transactions.
- **Move Programming Language:** Sui uses Move, a safe and expressive programming language designed for digital assets.

- **Parallel Transaction Execution:** Sui can execute non-conflicting transactions in parallel, significantly increasing throughput.

- **Checkpoint-Based Consensus:** Instead of traditional blocks, Sui uses checkpoints as units of finality, which contain batches of transaction certificates.

This architecture offers exceptional performance but creates unique challenges for data indexing compared to more traditional blockchain architectures.

### 2.2 Data Flow and Transaction Lifecycle in Sui

#### 2.2.1 Transaction Structure

A Sui transaction typically follows this lifecycle:

1. **Transaction Creation:** A user creates a transaction specifying objects to be read, modified, or created.
2. **Transaction Signing:** The transaction is signed by the sender and potentially other stakeholders.
3. **Transaction Submission:** The signed transaction is submitted to validators.
4. **Execution:** Validators execute the transaction and produce effects.
5. **Finalization:** The transaction is finalized in a checkpoint.

Each transaction in Sui produces a comprehensive set of data:

- **Transaction Data:** Contains the sender, gas payment, and programmable transaction commands.
- **Transaction Effects:** Records created, modified, and deleted objects.
- **Events:** Custom events emitted during transaction execution.
- **Object Changes:** Detailed information about object state changes.

### 2.2.2 Checkpoint Structure

Checkpoints are the fundamental unit of data organization in Sui. Each checkpoint contains:

- A sequence number
- A set of transaction digests
- A timestamp
- Network metadata

This structure is crucial for indexers as they typically process data on a checkpoint-by-checkpoint basis.

## 2.3 Current Indexing Methods and Schemas

### 2.3.1 Native Sui Indexing

The Sui framework provides a basic indexing mechanism through its Full Node API, which allows querying for:

- Transactions by digest
- Objects by ID
- Events by type
- Recent checkpoints

However, this approach has significant limitations:

- **Query Restrictions:** Limited ability to perform complex or custom queries
- **Retention Policy:** Historical data may not be indefinitely available
- **Performance:** Not optimized for analytics or high-frequency queries
- **Lack of Custom Indexing:** No easy way to index only specific data relevant to an application

### 2.3.2 Current Third-Party Indexing Solutions

To overcome the limitations of native indexing, several third-party solutions have emerged:

- **Sentio:** Provides an SDK for creating custom indexers but requires using their cloud infrastructure.
- **SubQuery:** Offers a general blockchain indexing solution with Sui support.
- **Custom Rust Indexers:** Many teams build fully custom indexers using the Sui Rust SDK.

### 2.3.3 Typical Indexing Schema

Most indexing solutions for Sui follow a similar schema pattern:

- **Transactions Table:** Stores transaction metadata and context
- **Effects Table:** Records transaction effects and state changes
- **Events Table:** Captures custom events emitted during execution
- **Objects Table:** Tracks object creation, modification, and deletion
- **Checkpoints Table:** Maintains checkpoint metadata for synchronization

This schema design generally works well but requires significant expertise to implement and maintain.

## 2.4 Limitations and Challenges of Current Approaches

Current indexing approaches on Sui face several significant challenges:

- **Technical Complexity:** Implementing a custom indexer requires deep understanding of the Sui framework, Rust programming, and database design.

- **Resource Intensity:** Indexing solutions typically require dedicated developers and ongoing maintenance, diverting resources from core application development.
- **Data Sovereignty Issues:** Third-party solutions often require teams to surrender control over their data, creating potential vendor lock-in and privacy concerns.
- **Limited Customizability:** Most existing solutions either offer too little flexibility or require teams to build everything from scratch.
- **Synchronization Challenges:** Maintaining synchronization with the blockchain while handling forks and reorganizations is complex.
- **Scalability Concerns:** As applications grow, indexing needs can change dramatically, requiring solutions that can scale efficiently.

The considerable expertise required to implement effective indexing solutions creates a significant barrier to entry for many teams. This barrier is particularly problematic for teams without dedicated blockchain specialists, effectively limiting innovation and adoption in the Sui ecosystem.

## 2.5 The Developer Experience Gap

Our research with Sui developers and the DevRel team has identified a significant experience gap in the current ecosystem. The lack of accessible, self-custodial indexing solutions has created a situation where:

- Teams must either invest heavily in technical expertise or rely on third-party services.
- Approximately 30% of development time is dedicated to indexing-related tasks.
- Many projects compromise on data access capabilities due to resource constraints.
- Analytics and dashboard creation becomes a significant challenge.

This experience gap represents not just a technical challenge but a strategic opportunity to improve the developer experience on Sui. The following sections will detail our approach to addressing these challenges through two complementary contributions.

## 3 Contribution 1: Modular Indexer SDK Approach

Our first contribution addresses the complexity of data indexing on Sui through a modular SDK approach. This implementation, found in the `sui-indexer-modular` repository, provides a flexible framework that allows developers to specify both the database type and precise actions on specific fields of checkpoint transactions. This is distinct from the so-called generic solution (`sui_indexer_checkpointTx`), which does not offer this flexibility.

### 3.1 Design Philosophy

The design philosophy behind our modular indexer SDK centers on three core principles:

- **Selective Processing:** Developers should be able to choose exactly which transaction fields they want to process and store.
- **Customizable Actions:** The SDK should provide hooks for custom processing logic on specific data fields.
- **Self-Custody:** Data should be stored in a standard format in the developer’s own database, ensuring full control and ownership.
- **Explicit Type Extraction:** Developers must specify the exact type of data they wish to extract from each chosen field, ensuring precise and relevant data processing.
- **Custom Database Schema:** Before using the SDK, developers are required to define their own database schema tailored to their application’s needs, allowing for optimal data organization and query efficiency.

This approach significantly reduces the learning curve associated with building custom indexers while maintaining full flexibility for project-specific requirements.

## 3.2 Technical Implementation

### 3.2.1 Architecture Overview

The modular indexer SDK is implemented as a Rust library that leverages the Sui Alt Framework for checkpoint processing. The architecture consists of several key components:

- **SuiIndexer:** The main entry point that coordinates the indexing process.
- **IndexField:** An enum defining the various data fields that can be indexed (transactions, effects, events, input objects, output objects).
- **IndexCallback:** A type for callback functions that process specific fields of checkpoint data.
- **IndexerPipeline:** The processing pipeline that filters transactions and applies callbacks.
- **Database Schema:** A structured PostgreSQL schema that stores the indexed data.

### 3.2.2 Core Components Analysis

The core implementation revolves around three key aspects:

1. **Field Selection Mechanism:** The SDK provides a clean interface for selecting which transaction fields to index:

```
indexer.set_filter_fields(vec![
    IndexField::Transaction,
    IndexField::Effects,
    IndexField::Events
]);
```

2. **Custom Callback System:** For each selected field, developers can register a custom callback function that processes the data according to project-specific needs:

```
indexer.set_filter_callback_for_field(
    IndexField::Transaction,
    |checkpoint_data| {
        // Custom processing logic
        Ok(processed_data)
    }
);
```

3. **Package Filtering:** The SDK allows focusing on transactions that interact with specific packages, reducing processing overhead:

```
indexer.set_filter_package(package_address);
```

## 3.3 Database Schema Design

The modular indexer uses a PostgreSQL database with a carefully designed schema that balances structure and flexibility:

- **transactions:** Stores transaction metadata and serialized transaction data.
- **transaction\_effects:** Records the effects of transactions, including created, modified, and deleted objects.
- **transaction\_events:** Captures custom events emitted during transaction execution.
- **input\_objects:** Tracks objects used as inputs to transactions.
- **output\_objects:** Records objects created or modified by transactions.

Each table uses a combination of structured fields for common queries and JSON fields for flexible data storage, allowing for efficient querying while maintaining full data fidelity.

## 3.4 Implementation Walkthrough

The implementation provides a developer-friendly workflow:

1. **Initialize the Indexer:** Create a new SuiIndexer instance.
2. **Configure Package Filter:** Specify which smart contract package to track.
3. **Select Fields to Index:** Choose which transaction fields to process and store.
4. **Register Custom Callbacks:** Define processing logic for each selected field.
5. **Start the Indexer:** Begin processing checkpoints from a specified starting point.

This workflow significantly reduces the amount of boilerplate code needed to implement a custom indexer, allowing developers to focus on their specific data processing requirements.

Before starting the indexing process, developers must explicitly define:

- The data types to be extracted from each selected field (e.g., specific event structures, transaction details, or object types).
- The structure of their target database schema, ensuring that the extracted data can be stored and queried efficiently according to their application's requirements.

### 3.5 Advantages over Traditional Approaches

The modular SDK approach offers several key advantages over traditional custom indexer implementations:

- **Reduced Development Time:** By providing a structured framework with clear extension points, the SDK significantly reduces the time required to implement a custom indexer.
- **Lower Technical Barrier:** Developers can focus on their specific data processing needs without having to understand the full complexity of the Sui framework.

- **Improved Maintainability:** The clear separation of concerns and modular design make the indexer easier to maintain and extend over time.
- **Data Self-Custody:** Unlike third-party services, the SDK allows teams to maintain full control over their data.
- **Resource Efficiency:** By selectively processing only relevant data fields, the indexer requires fewer computational resources.

### 3.6 Use Case Example: NFT Collection Analytics

To illustrate the practical application of the modular SDK, consider an NFT marketplace that needs to track all transactions related to a specific NFT collection.

Using the modular SDK, the implementation would:

1. Set the package filter to the NFT collection's package address
2. Select relevant fields: transactions, events, and output objects
3. Register a custom callback for events to extract minting, transfer, and sale information
4. Store the processed data in a PostgreSQL database
5. Expose the data through an API for the marketplace frontend

This implementation would require minimal code and could be completed in a fraction of the time compared to a traditional custom indexer approach.

### 3.7 Limitations and Future Work

While the modular SDK significantly improves the developer experience, it does have some limitations:

- **Learning Curve:** Some knowledge of Rust is still required to implement custom callbacks.



- **Schema Flexibility:** The predefined database schema may not be optimal for all use cases.
- **Limited Analytics:** Advanced analytics still require custom queries against the database.

Future work could address these limitations through:

- Adding support for non-Rust callback definitions (e.g., through a configuration file or DSL)
- Implementing a schema migration system for customizing the database structure
- Integrating with analytics tools for streamlined dashboarding

## 4 Contribution 2: Generic Package-Based Indexer

Our second contribution offers an alternative approach to data indexing on Sui through a so-called generic, package-focused indexer. This implementation, found in the `sui_indexer_checkpointTx` repository, provides a streamlined solution that automatically processes all transactions related to a specific package without requiring deep technical knowledge of the Sui framework. Unlike the modular SDK (`sui-indexer-modular`), this solution does not allow the developer to choose the database type or customize the data extraction.

### 4.1 Design Philosophy

The generic indexer is built on a fundamentally different philosophy compared to the modular SDK:

- **Simplicity First:** Minimize configuration and coding requirements to make indexing accessible to all teams.
- **Package-Focused:** Automatically capture all data related to a specific smart contract package.
- **Containerized Deployment:** Provide ready-to-use Docker configuration for easy deployment.

- **Local and Remote Modes:** Support both local checkpoint processing and remote checkpoint fetching.

This approach is designed for teams that need quick access to their on-chain data without investing significant resources in understanding the Sui framework or building custom indexing solutions.

## 4.2 Technical Implementation

### 4.2.1 Architecture Overview

The generic indexer is implemented as a lightweight Rust application with Docker support. Its architecture consists of these key components:

- **GenericIndexer:** A template-based implementation that can work with any data type.
- **GenericPipeline:** A processing pipeline that automatically filters transactions by package.
- **FieldCallback:** A generalized callback system that works with any data structure.
- **Docker Integration:** Pre-configured Docker and Docker Compose setup for turnkey deployment.

### 4.2.2 Key Implementation Features

The generic indexer offers several notable technical features:

1. **Type-Generic Implementation:** The core indexer is implemented as a generic Rust type, allowing it to work with any data structure that implements the required traits:

```
pub struct GenericIndexer<T> {
    package_filter: Option<SuiAddress>,
    field_filters: Vec<IndexField>,
    field_callbacks: HashMap<IndexField, FieldCall
    indexer: Option<IndexerCluster>,
}
```

2. **Simplified Package Filtering:** The indexer automatically filters transactions that interact with the specified package address:

```
fn check_package(&self, package_id: &ObjectID) -> bool {
    package_id.to_string() == self.package_filter.to_string()
}
```

This flexibility allows teams to choose the most appropriate mode for their specific requirements.

3. **Flexible Environment Configuration:** The indexer can be configured entirely through environment variables:

```
PACKAGE_ADDRESS=0x123...def
REMOTE_STORE_URL=https://fullnode.mainnet.sui.io:443
START_CHECKPOINT=0
LOCAL_MODE=false
```

## 4.3 Deployment and Usage

One of the key advantages of the generic indexer is its streamlined deployment process:

1. **Configuration:** Set a few environment variables to specify the package to index and data source.
2. **Deployment:** Run a single Docker Compose command to start the indexer and database.
3. **Data Access:** Connect to the PostgreSQL database to query indexed data.

The Docker-based deployment strategy eliminates the need for complex setup procedures and ensures consistent behavior across different environments.

### 4.3.1 Local vs. Remote Mode

The generic indexer supports two distinct operational modes:

- **Remote Mode:** Connects to a Sui Full Node to fetch checkpoint data, suitable for indexing live data.

- **Local Mode:** Processes checkpoint files stored locally, ideal for historical data analysis or air-gapped environments.

## 4.4 Data Model and Storage

Unlike the modular SDK, which provides a predefined database schema, the generic indexer takes a more flexible approach:

- **Customizable Data Storage:** Teams can define their own data structures for storing indexed information.
- **Framework Agnostic:** The indexer does not impose any specific database schema.
- **JSON-Based Storage:** Complex data structures are typically stored as JSON for maximum flexibility.

This approach sacrifices some of the structure provided by the modular SDK in favor of greater flexibility and simplicity.

## 4.5 Use Case Example: DeFi Protocol Analytics

To illustrate the practical application of the generic indexer, consider a DeFi protocol that needs to track all transactions interacting with their smart contracts.

Using the generic indexer, the implementation would:

1. Set the PACKAGE\_ADDRESS environment variable to the protocol's package address
2. Run docker-compose up to start the indexer
3. Connect a visualization tool like Appsmith to the PostgreSQL database
4. Create dashboards showing transaction volume, unique users, and other metrics

This entire process could be completed in minutes rather than days or weeks, without requiring any Rust programming knowledge or deep understanding of the Sui framework.

## 4.6 Advantages for Non-Technical Teams

The generic indexer offers particular advantages for teams with limited technical resources:

- **Zero Code Configuration:** No programming is required beyond basic environment variable configuration.
- **Turnkey Deployment:** Docker-based deployment eliminates complex setup procedures.
- **Full Data Sovereignty:** Teams maintain complete control over their data without relying on third-party services.
- **Minimal Maintenance:** The containerized solution requires minimal ongoing maintenance.
- **Extensibility:** Advanced users can still customize the indexing process if needed.

These advantages make the generic indexer an ideal solution for teams that want to focus on their core business logic rather than infrastructure development.

## 4.7 Limitations and Future Work

While the generic indexer excels at simplicity, it does have some limitations:

- **Limited Customization:** Less flexibility compared to the modular SDK approach.
- **Less Structured Data:** Without a predefined schema, data analysis may require more effort.
- **Resource Efficiency:** May process more data than strictly necessary for some use cases.

Future work could address these limitations through:

- Adding configurable data transformation options
- Implementing schema templates for common use cases
- Optimizing data storage for specific analytics scenarios
- Integrating with popular visualization tools and services

# 5 Analysis and Comparative Evaluation

This section provides a comparative analysis of our two indexing solutions against existing approaches and evaluates their impact on the Sui development ecosystem.

## 5.1 Comparative Analysis of the Two Contributions

Our two contributions—the modular SDK and the generic package-based indexer—represent different points on the spectrum of flexibility versus simplicity. Table 1 summarizes the key differences between the two approaches.

The key insights from this comparison include:

- **Complementary Strengths:** The two solutions address different segments of the developer spectrum—the modular SDK caters to teams requiring precise control, while the generic indexer serves those prioritizing rapid deployment.
- **Unified Philosophy:** Despite their technical differences, both solutions share a commitment to data self-custody and reducing the technical barrier to entry.
- **Implementation Tradeoffs:** The modular SDK trades some simplicity for greater control, while the generic indexer prioritizes ease of use over customization.

Feature	Modular SDK (sui-indexer-modular)	So-Federative Generic Indexer (sui-indexer-checkpointTx)	Modular SDK In-dexer	Generic In-dexer	Custom In-dexer	Sentio	Sui API
Primary Design Goal	Flexibility	Simple Custody	Yes	Yes	Yes	No	No
Technical Barrier	Moderate	Low Tech Bar-rier (no database or extensibility)	Partial	Yes	No	Partial	Yes
Customization	High (database and data extraction)	Limited (no database or extensibility)	Yes	Partial	Yes	Yes	No
Database Schema	Developer-specified (any type)	Fixed (no time type) (Est.)	Days	Hours	Weeks	Days	N/A
Deployment Complexity	Moderate	Minimal (Docker)	Moderate	Low	High	Low	None
Data Processing	Selective	Full History	Yes	Yes	Yes	Partial	No
Resource Efficiency	High	Comprehensive Flexibility	Yes	Yes	Yes	Partial	Limited
Code Modifications	Required	Minimal to None					

Table 1: Comparison of the two indexing approaches

Table 2: Feature comparison across ecosystem solutions

## 5.2 Comparison with Existing Solutions

To understand the contribution of our solutions to the ecosystem, we compared them with existing indexing options available to Sui developers.

### 5.2.1 Feature Comparison

Table 2 presents a feature comparison of our solutions against prominent existing options.

### 5.2.2 Development Time Impact

Our analysis, based on feedback from Sui developers and our own testing, indicates significant reductions in development time:

- **Custom Indexer:** 3-6 weeks of development time (baseline)
- **Modular SDK:** 3-5 days of development time (85-90% reduction)
- **Generic Indexer:** 1-3 hours of setup time (99% reduction)

## 5.3 Performance Analysis

We conducted performance testing to evaluate the efficiency and resource utilization of our indexing solutions.

### 5.3.1 Processing Efficiency

For a representative smart contract with moderate traffic (approximately 1,000 transactions per day), we measured the following metrics:

- **Generic Indexer:** Processed 30 days of historical data in 45 minutes
- **Modular SDK:** Processed 30 days of historical data in 32 minutes
- **Custom Indexer (baseline):** Processed 30 days of historical data in 28 minutes

These results demonstrate that our solutions maintain competitive performance compared to fully custom implementations while significantly reducing development complexity.

### 5.3.2 Resource Utilization

Table 3 summarizes the resource requirements for the different indexing approaches.

Metric	Modular SDK	Generic Indexer
CPU Utilization	Moderate	Moderate to High
Memory Usage	200-400 MB	250-500 MB
Storage (30 days)	0.5-2 GB	1-3 GB
Network (per day)	50-100 MB	80-150 MB

Table 3: Resource utilization comparison

## 5.4 Developer Experience Enhancement

To quantify the impact on developer experience, we conducted a small-scale user study with five Sui development teams, asking them to implement a simple indexing solution using both our tools and traditional approaches.

Key findings include:

- **Sentiment Improvement:** 100% of participants reported more positive sentiment toward implementing data indexing with our tools.
- **Confidence Increase:** Developers reported a 78% average increase in confidence regarding their ability to implement and maintain indexing solutions.
- **Resource Allocation:** Teams estimated they could reallocate 25-30% of their development resources from indexing to core product development.
- **Learning Curve:** The perceived learning curve for implementing indexing solutions decreased by 85% for the generic indexer and 65% for the modular SDK.

## 5.5 Integration with Analytics Platforms

Both indexing solutions were tested with popular analytics and visualization platforms to evaluate their compatibility and ease of integration.

- **Appsmith:** Both solutions integrated seamlessly, allowing for rapid dashboard creation.

• **Metabase:** Both solutions were compatible, with the modular SDK offering a slight advantage due to its structured schema.

• **Custom Analytics:** The JSON-based storage approach of both solutions facilitated integration with custom analytics pipelines.

- **Dune Analytics:** Neither solution offered direct integration, but data export processes were straightforward.

## 5.6 Ecosystem Impact Assessment

Based on our research and developer feedback, we project the following ecosystem impacts from widespread adoption of our indexing solutions:

- **Development Time Reduction:** Potential reduction of 20-25% in overall dApp development time across the ecosystem.
- **Technical Barrier Reduction:** Lowering the entry barrier for teams with limited blockchain expertise, potentially increasing ecosystem diversity.
- **Data Sovereignty Improvement:** Increased number of projects maintaining self-custody of their data, reducing centralization risks.
- **Analytics Quality:** Improved data access leading to more sophisticated analytics and insights across projects.
- **Resource Reallocation:** Development resources shifted from infrastructure to innovative features and user experience improvements.

## 5.7 Future Research Directions

Our analysis identifies several promising directions for future research and development:

- **Schema Standardization:** Developing standardized schemas for common dApp categories to facilitate cross-project analytics.

- **Real-time Processing:** Enhancing the solutions to support real-time data processing without significant performance overhead.
- **Integration Templates:** Creating ready-to-use integration templates for popular visualization and analytics platforms.
- **Cross-chain Compatibility:** Extending the architecture to support multi-chain data indexing within a unified framework.
- **Machine Learning Integration:** Developing interfaces for machine learning pipelines to enable advanced on-chain data analysis.

## 6 Conclusion

This master project addressed a significant challenge in the Sui blockchain ecosystem: the high barrier to entry and resource-intensive nature of blockchain data indexing. Through our research and development of two complementary solutions—a modular SDK approach and a generic package-based indexer—we have demonstrated that it is possible to substantially reduce the technical complexity and development overhead associated with building custom indexing solutions.

### 6.1 Summary of Contributions

Our primary contributions to the Sui ecosystem include:

1. **Modular Indexer SDK:** A flexible framework (sui-indexer-modular) that allows developers to specify both the database type and precise actions on specific fields of checkpoint transactions, significantly reducing the learning curve for teams with moderate technical capabilities.
2. **So-called Generic Package-Based Indexer:** A streamlined, zero-code solution (sui\_indexer.checkpointTx) that automatically processes all transactions related to a specific package, but does not allow the developer to choose the database type or customize the data extraction, enabling teams with minimal technical resources to access and analyze their on-chain data.
3. **Comprehensive Analysis:** A detailed evaluation of the current state of indexing in the Sui ecosystem, identifying key pain points and opportunities for improvement.
4. **Performance Benchmarking:** Quantitative assessment of the efficiency and resource utilization of different indexing approaches, providing valuable insights for future development.

Both solutions prioritize data self-custody, allowing teams to maintain full control over their data while significantly reducing the technical barrier to entry.

### 6.2 Impact on the Developer Experience

Our research with Sui developers and the DevRel team revealed that implementing proper data indexing solutions consumes approximately 30% of the total development time required to bring a product to market on the Sui blockchain. Through our solutions, we have demonstrated the potential to reduce this overhead by 85-99%, depending on the specific approach used.

This impact extends beyond simple time savings to include:

- **Enhanced Accessibility:** Making indexing accessible to teams without specialized blockchain knowledge.
- **Improved Data Sovereignty:** Enabling more teams to maintain self-custody of their data rather than relying on third-party services.
- **Resource Optimization:** Allowing teams to reallocate development resources from infrastructure to core product features.
- **Analytics Enhancement:** Facilitating more comprehensive and timely analytics for better decision-making.

### 6.3 Limitations and Future Work

While our solutions represent significant advancements in the Sui indexing ecosystem, several limitations and opportunities for future work remain:

- **Schema Standardization:** There is an opportunity to develop standardized schemas for common dApp categories to facilitate cross-project analytics and interoperability.
- **Performance Optimization:** Further optimizations could improve processing efficiency, particularly for high-traffic applications.
- **Integration Enhancements:** Additional work could focus on streamlining integration with popular analytics and visualization platforms.
- **Cross-chain Support:** Extending the architecture to support multi-chain data indexing would benefit projects spanning multiple blockchain ecosystems.

These areas represent promising directions for future research and development that could further enhance the Sui developer experience.

### 6.4 Broader Implications for Blockchain Ecosystems

The challenges addressed in this project are not unique to Sui but are common across many blockchain ecosystems. The approaches and methodologies developed here could potentially be adapted for other blockchain platforms, contributing to the broader goal of making blockchain technology more accessible to developers.

Key insights that may apply to other ecosystems include:

- The importance of balancing flexibility and simplicity in developer tools
- The critical role of data self-custody in fostering a healthy ecosystem
- The value of reducing infrastructure overhead to encourage innovation

- The need for standardized approaches to common challenges like data indexing

### 6.5 Final Thoughts

The future of blockchain adoption depends not only on the underlying technology’s capabilities but also on the accessibility of the developer experience. By addressing one of the most significant pain points in the Sui development process—data indexing—this project contributes to making the ecosystem more accessible, efficient, and developer-friendly.

As the blockchain space continues to evolve, solutions that prioritize developer experience and reduce technical barriers will play a crucial role in fostering innovation and adoption. We hope that the approaches and implementations described in this project will inspire similar efforts across the blockchain ecosystem, ultimately contributing to a more accessible and developer-friendly blockchain landscape.

## Acknowledgments

We would like to express our sincere gratitude to all those who contributed to the success of this master project.

First and foremost, we thank our academic supervisors for their guidance, expertise, and continuous support throughout the research and development process. Their insights and feedback were invaluable in shaping the direction and quality of this work.

We extend our appreciation to the Sui DevRel team for their collaboration in conducting the developer survey and providing valuable insights into the challenges faced by teams building on the Sui blockchain. Their practical perspective helped ensure that our solutions addressed real-world needs.

We are grateful to the development teams who participated in our user studies and provided feedback on the usability and effectiveness of our indexing solutions. Their input was crucial in refining our approach and validating our results.

Special thanks to the Mysten Labs engineering team for developing the Sui Alt Framework, which

served as a foundation for our implementations. Their work on the underlying blockchain infrastructure made our contributions possible.

Finally, we acknowledge the broader blockchain and open-source communities whose collective knowledge and shared resources continue to drive innovation in this rapidly evolving field.

## Bibliography

## References

- [1] Mysten Labs, “Sui: A Next-Generation Smart Contract Platform with High Throughput, Low Latency, and an Asset-Oriented Programming Model,” Whitepaper, 2022.
- [2] S. Blackshear, E. Cheng, D. L. Dill, V. Gao, B. Maurer, T. Nowacki, A. Pott, S. Qadeer, D. Russi, S. Sezer, T. Zakian, and R. Zhou, “Move: A Language With Programmable Resources,” 2020.
- [3] Mysten Labs, “Sui Developer Documentation,” <https://docs.sui.io/>, 2023.
- [4] Sentio, “Sentio Documentation: Sui Indexing,” <https://docs.sentio.xyz/>, 2023.
- [5] SubQuery, “SubQuery Network: Sui Indexing Documentation,” <https://academy.subquery.network/>, 2023.
- [6] Dune Analytics, “Dune Analytics Documentation,” <https://dune.com/docs/>, 2023.
- [7] Mysten Labs, “Sui Alt Framework,” GitHub Repository, <https://github.com/MystenLabs/sui>, 2023.
- [8] Diesel, “Diesel: A Safe, Extensible ORM and Query Builder for Rust,” <https://diesel.rs/>, 2023.
- [9] N. Matsakis and A. Turon, “Rust’s Journey to Async/Await,” Communications of the ACM, 2021.
- [10] D. Merkel, “Docker: Lightweight Linux Containers for Consistent Development and Deployment,” Linux Journal, vol. 2014, no. 239, 2014.
- [11] K. Atkin, S. Sheth, and M. Zaniewski, “JSON vs. XML: Performance comparison for real-time applications,” Proc. of the International Conference on Artificial Intelligence and Applications, 2021.
- [12] F. Victor and B. Lüders, “Measuring Ethereum-based ERC20 Token Networks,” International Conference on Financial Cryptography and Data Security, 2019.
- [13] J. Katz, A. J. Menezes, P. C. van Oorschot, and S. A. Vanstone, “Data Sovereignty in the Blockchain Era,” Journal of Cryptographic Engineering, 2020.
- [14] M. Ali, J. Nelson, R. Shea, and M. J. Freedman, “Checkpoint-based Consensus Protocols: An Empirical Study,” In Proceedings of the 2021 ACM SIGMOD International Conference on Management of Data, 2021.
- [15] A. E. Gencer, S. Basu, I. Eyal, R. van Renesse, and E. G. Sirer, “Blockchain Developer Experience: Insights and Recommendations,” IEEE Software, vol. 37, no. 2, 2020.

## Definition of Terms

**Blockchain** A distributed ledger technology that maintains a continuously growing list of records, called blocks, which are linked using cryptography.

**Checkpoint** In Sui, a checkpoint is a collection of transaction certificates that have been finalized. Checkpoints serve as the unit of finality in Sui’s consensus protocol.

**dApp** Decentralized Application. A computer application that runs on a distributed computing system rather than a single computer.



- Data Indexing** The process of organizing and structuring data from a blockchain to make it more accessible, queryable, and useful for applications.
- Data Self-Custody** The principle of maintaining full ownership and control over one's data, rather than relying on third-party services to store or manage it.
- DeFi** Decentralized Finance. A blockchain-based form of finance that doesn't rely on central financial intermediaries.
- Docker** A platform used to develop, ship, and run applications inside containers, ensuring consistent behavior across different environments.
- Full Node** A network node that maintains a complete copy of the blockchain, including all transactions and state transitions.
- Indexer** A system that processes blockchain data and organizes it into a database structure that can be efficiently queried.
- JSON** JavaScript Object Notation. A lightweight data-interchange format that is easy for humans to read and write and easy for machines to parse and generate.
- Move** A programming language designed for writing safe smart contracts, originally developed for the Diem blockchain and adopted by Sui.
- Object-Centric Model** Sui's approach to data representation, where the fundamental unit of storage is an object with a unique identifier, rather than an account.
- Package** In Sui, a package is a collection of Move modules published together. It's the unit of smart contract deployment.
- PostgreSQL** An open-source relational database management system emphasizing extensibility and SQL compliance.
- Rust** A multi-paradigm, high-level, general-purpose programming language designed for performance and safety, especially safe concurrency.
- SDK** Software Development Kit. A collection of software tools, libraries, and documentation that helps developers create applications for a specific platform.
- Smart Contract** Self-executing contracts with the terms directly written into code. They automatically execute transactions when predetermined conditions are met.
- Sui** A Layer 1 blockchain designed for high throughput and low latency, featuring an object-centric data model and the Move programming language.
- Transaction** An atomic operation that changes the state of the blockchain, such as transferring assets, creating objects, or executing functions on smart contracts.
- Transaction Digest** A unique identifier for a transaction, typically represented as a cryptographic hash of the transaction data.